

SUMMARY

Data Engineer with 5+ years building production-grade ETL/ELT pipelines on AWS and Azure. Delivered event-driven pipelines, cutting runtime by 30% and end-to-end observability, reducing incident response by 35%. Focused on Gen AI/RAG systems, data quality, and IAM least-privilege security.

SKILLS

Cloud & AWS

Data Engineering

Databases

AI & ML

Programming & DevOps

Data Governance

EC2, S3, EMR, Lambda, Glue, Redshift, Athena, Kinesis, SQS, SNS, Step Functions, CloudWatch

PySpark, Apache Spark, Kafka, Apache Iceberg, Airflow, dbt, ETL/ELT, CDC, Hadoop, SSIS, Spark Structured Streaming

PostgreSQL, DynamoDB, Snowflake, SQL, NoSQL, Netezza

LangChain, RAG, LLM, HuggingFace, OpenAI, ChromaDB, scikit-learn, pandas, NumPy, MLflow

Python, Scala, Bash, Flask, Django, Terraform, Docker, Kubernetes, CI/CD, IAM, RBAC, Spring Boot

Data Quality, Data Lineage, Data Observability, Data Modeling, Splunk, JIRA, Agile

EXPERIENCE

AWS Data Engineer · State Farm, Chicago IL

Jul 2024 – Present

Designed scalable ETL pipelines using AWS Glue, PySpark, and Python integrating SQL/NoSQL and cloud storage, cutting data processing time by 25%.

Optimized production pipelines with EMR, PySpark, Scala ingesting from S3 and Apache Iceberg; reduced runtime by 30% via improved partitioning and parallel execution.

Managed cloud infrastructure with Terraform, Python, and SSIS for data integration across Hadoop and Netezza environments, reducing infrastructure costs by 30%.

Implemented end-to-end observability with Splunk, reducing incident response time by 35% across business-critical insurance pipelines.

Enforced least-privilege IAM policies across Lambda, Glue, API Gateway, and cross-account S3; configured STS role assumption for CI/CD to eliminate long-lived credentials.

Improved analytical query performance by 20% using Redshift Spectrum and Starburst as federated query engines over Apache Iceberg tables in S3.

Built data quality monitoring with CloudWatch and custom Python scripts, reducing pipeline errors by 25% across 10+ daily scheduled runs.

Data Engineer · Bank of America, Hartford CT

Aug 2022 – Jun 2024

Led migration of large-scale data platform to Delta Lake on Databricks using AWS Glue, Python, and Scala; achieved 40% improvement in processing efficiency.

Designed Snowflake pipelines with Streams, Tasks, and dynamic data masking; improved query performance by 25% with RBAC for compliance datasets.

Built production ETL pipelines on AWS (Glue, Lambda, Step Functions) supporting 15+ daily financial data runs.

Optimized Spark and SQL jobs with partition pruning, broadcast joins, and predicate pushdown, reducing the average runtime of the job by 35%.

Redesigned IAM policies and implemented the assumption of STS role in Jenkins CI/CD, reducing the model deployment time by 40%.

Authored dbt models to centralize transformation logic for finance/compliance reporting with automated Python validation and CloudWatch monitoring.

Python Developer · Dentsu, Hyderabad India

Aug 2020 – Dec 2021

Developed scalable ETL pipelines using Python, PySpark, and Azure Data Factory, improving data processing efficiency by 30%.

Orchestrated multi-step workflows with Apache Airflow DAGs, automating dependency management and reducing manual intervention.

Built REST APIs with Flask and Django enabling self-service data access for business teams, significantly reducing ad-hoc engineering requests.

Implemented batch processing with Python and Spring Batch; automated CI/CD deployments, improving processing speed by 30%.

EDUCATION

New England College, Henniker NH

Jan 2022 – May 2023

M.S. in Computer and Information Systems

GPA: 3.95

CVR College of Engineering, India

Aug 2016 – Apr 2020

B.Tech in Electronics and Communication Engineering

GPA: 3.56

CERTIFICATIONS

AWS Certified Cloud Practitioner — Issued 2023

AWS Certified Developer, Associate — Issued 2025